

# Configuration Manual for Cloud-Based Time Series Prediction of CPU Load Through Advanced Deep Learning Models

MSc Research Project  
Masters In Cloud Computing

Rohan Ajila  
Student ID: 22249729

School of Computing  
National College of Ireland

Supervisor: Prof. Vikas Sahni

**National College of Ireland**  
**MSc Project Submission Sheet**  
**School of Computing**



**Student Name:** Rohan Ajila  
**Student ID:** 22249729  
**Programme:** MSc. Cloud Computing **Year:** 2024-25  
**Module:** MSc. Research Project  
**Lecturer:** Prof. Vikas Sahni  
**Submission Due Date:** 24 April 2025  
**Project Title:** Cloud-Based Time Series Prediction of CPU Load Through Advanced Deep Learning Models  
**Word Count:** 1410 **Page Count:** 20

I hereby certify that the information contained in this (my submission) is information pertaining to research I conducted for this project. All information other than my own contribution will be fully referenced and listed in the relevant bibliography section at the rear of the project.

ALL internet material must be referenced in the bibliography section. Students are required to use the Referencing Standard specified in the report template. To use other author's written or electronic work is illegal (plagiarism) and may result in disciplinary action.

**Signature:** Rohan Ajila  
**Date:** 24 April 2025

**PLEASE READ THE FOLLOWING INSTRUCTIONS AND CHECKLIST**

Attach a completed copy of this sheet to each project (including multiple copies)	<input type="checkbox"/>
<b>Attach a Moodle submission receipt of the online project submission,</b> to each project (including multiple copies).	<input type="checkbox"/>
<b>You must ensure that you retain a HARD COPY of the project,</b> both for your own reference and in case a project is lost or mislaid. It is not sufficient to keep a copy on computer.	<input type="checkbox"/>

Assignments that are submitted to the Programme Coordinator Office must be placed into the assignment box located outside the office.

<b>Office Use Only</b>	
Signature:	
Date:	
Penalty Applied (if applicable):	

# Configuration Manual for Cloud-Based Time Series Prediction of CPU Load Through Advanced Deep Learning Models

Rohan Ajila  
22249729

## 1 Introduction

This configuration manual aims to provide a step-by-step guide for setting up the cloud-native environment for the development and code implementation for the research project titled “Cloud-Based Time Series Prediction of CPU Load Through Advanced Deep Learning Models”. The main goal of the research is to forecast CPU utilization in Microsoft Azure virtual machines using historical time series data, which helps in more proactive resource scaling in cloud environments. The study compares and evaluates the efficiency of different traditional models like Decision Tree and Gradient Boosting to the deep learning models like LSTM, BiLSTM and Multihead BiLSTM for time series forecasting.

The cloud-based implementation in this project makes use of the Amazon Web Services (AWS) like EC2 for computational purposes, S3 for storage of data and cloud 9 as the integrated development environment. Time series data from azure log files is processed using Python tools and libraries and the performance metrics used to evaluate the performance of the models are Root Mean Squared Error (RMSE) and  $R^2$  Score.

The following manual is presented as follows: Section 2 discusses the system specifications setup on the cloud for the implementation. Section 3 displays the software installation, AWS environment setup and the Python Libraries required. Section 4 shows the implementation and evaluation of the forecasting models. Section 5 provides the conclusions of the research and section 6 lists all the references used in this project.

## 2 Cloud Environment System Specification

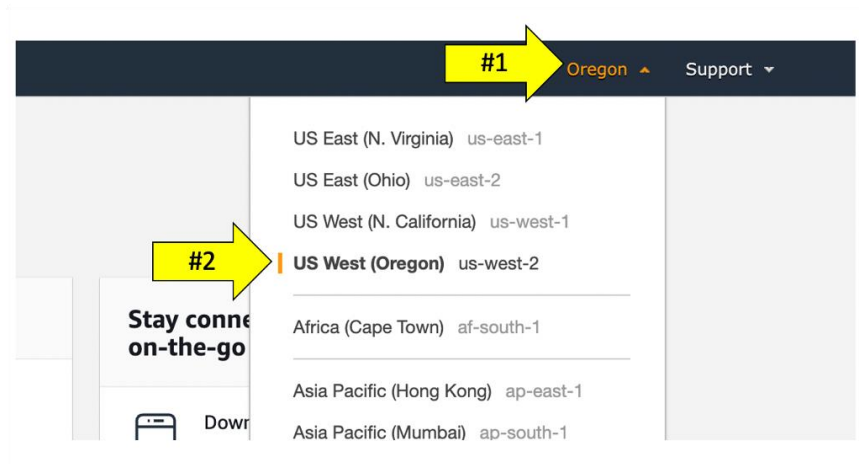
### 2.1 Prerequisite:

The user should have a basic knowledge of python syntax and working flow of python > 3.10 or later version, and also the concepts like machine learning and deep learning. They should also be familiar with the EC2 instances, Cloud9, S3 Bucket.

### 2.2 Cloud 9 Setup:

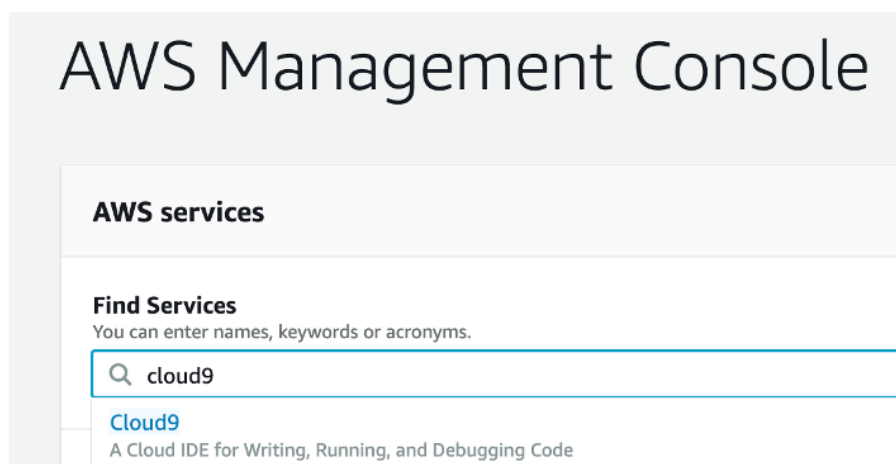
Set up the AWS Cloud9 IDE on an EC2 instance for development and project Implementation (Jupyter Notebook).

STEP 1: Open the AWS management console and select the desired location.



**Figure 1 : Region selection on AWS Console**

STEP 2: Find Cloud 9 in the search panel of the AWS Management Console



**Figure 2 : Cloud 9 on AWS Console**

STEP 3: Create an environment and setup the environment as shown in figure 3. Make sure to select create new EC2 instance for environment using direct access.

## Environment settings

### Environment type [Info](#)

Run your environment in a new EC2 instance or an existing server. With EC2 instances, you can connect directly through Secure Shell (SSH) or connect via AWS Systems Manager (without opening inbound ports).

- ☒ Create a new EC2 instance for environment (direct access)  
Launch a new instance in this region that your environment can access directly via SSH.
- ☐ Create and run in remote in remote server (SSH connection)  
Configure the secure connection to the remote server for your environment.

### Instance type

- ☐ t2.micro (1 GiB RAM + 1 vCPU)  
Free-tier eligible. Ideal for educational users and exploration.
- ☐ t3.small (2 GiB RAM + 2 vCPU)  
Recommended for small-sized web projects.
- ☐ m5.large (8 GiB RAM + 2 vCPU)  
Recommended for production and general-purpose development.
- ☒ Other instance type  
Select an instance type.

t3.medium

### Platform

- ☒ Amazon Linux
- ☐ Ubuntu Server 18.04 LTS

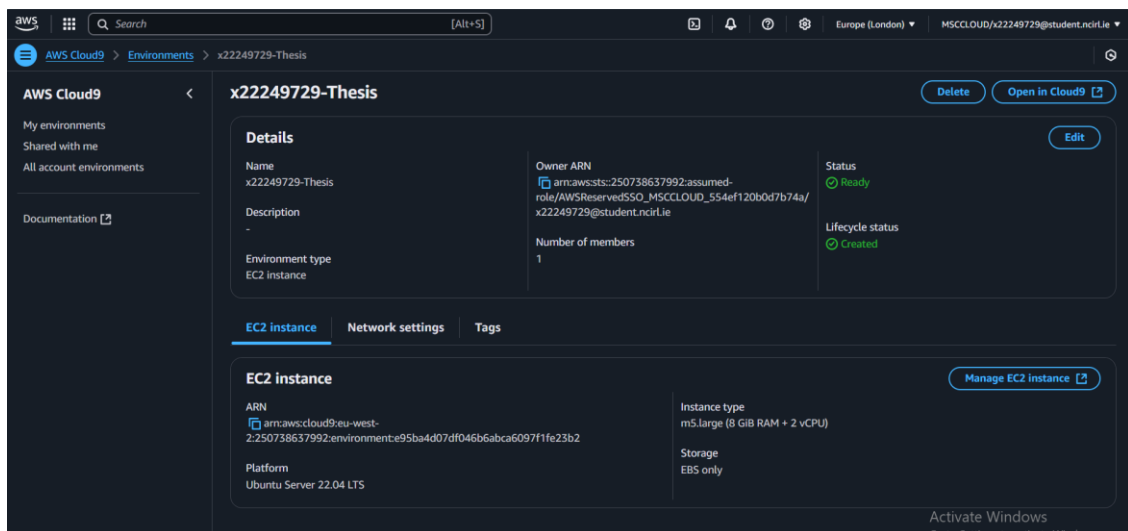
### Cost-saving setting

Choose a predetermined amount of time to auto-hibernate your environment and prevent unnecessary charges. We recommend a hibernation settings of half an hour of no activity to maximize savings.

After 30 minutes (default)

### IAM role

**Figure 3 : Environment settings**



**Figure 4: Cloud 9 environment setup**

STEP 4: Cloud 9 environment is now setup as shown in Figure 4.

## 2.3 EC2 Configuration:

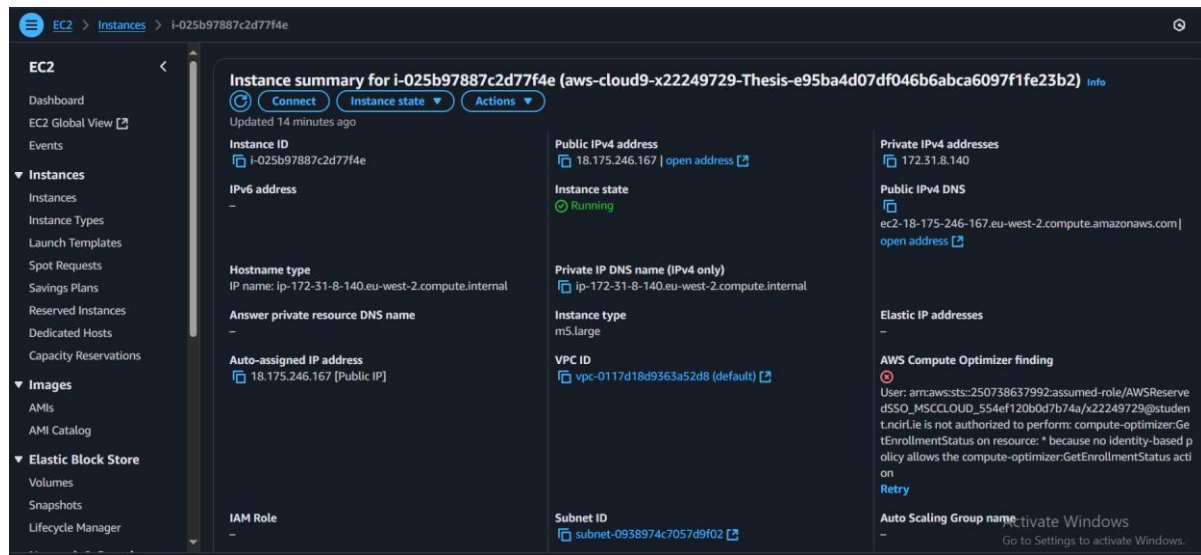


Figure 5 : EC2 instance connected to the Cloud 9 Environment

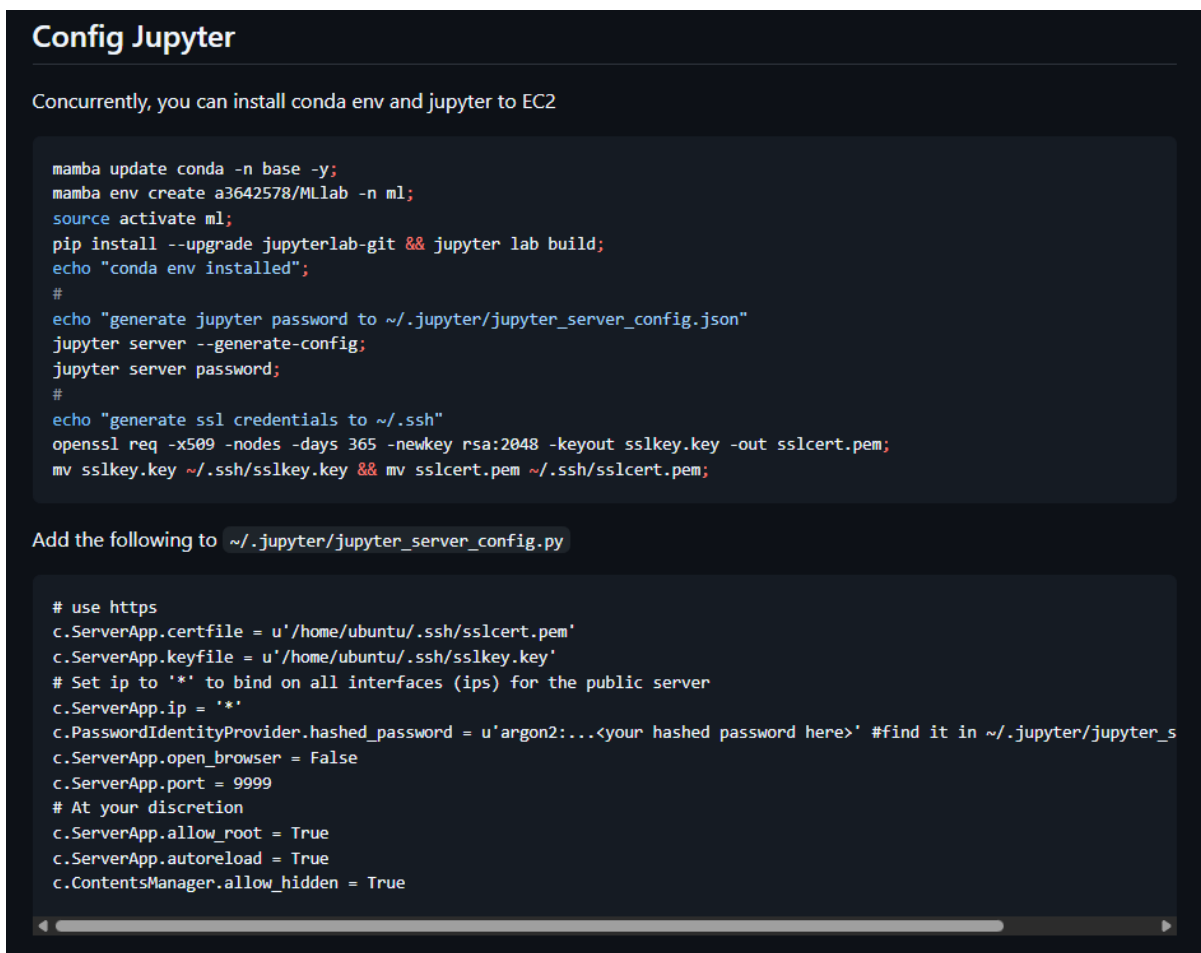


Figure 6 : Configure Jupyter notebook on EC2

## Launch Jupyter from Cloud9

Set Cloud9 Project Preference Stop my environemnt to never Now launch jupyter from Cloud 9's terminal

## Setup tunneling from local

Ubuntu / MacOS:

```
ssh -i private_key_to_ec2.pem -Nfl localhost:9999:localhost:9999 ubuntu@ec2_host
```

Windows: Using MobaXterm

Figure 7 : Launch Jupyter notebook from Cloud 9

STEP 1: The EC2 instance needs to be setup to run Jupyter notebook, follow the figure 6 and 7 to setup to run Jupyter directly from cloud 9 IDE.

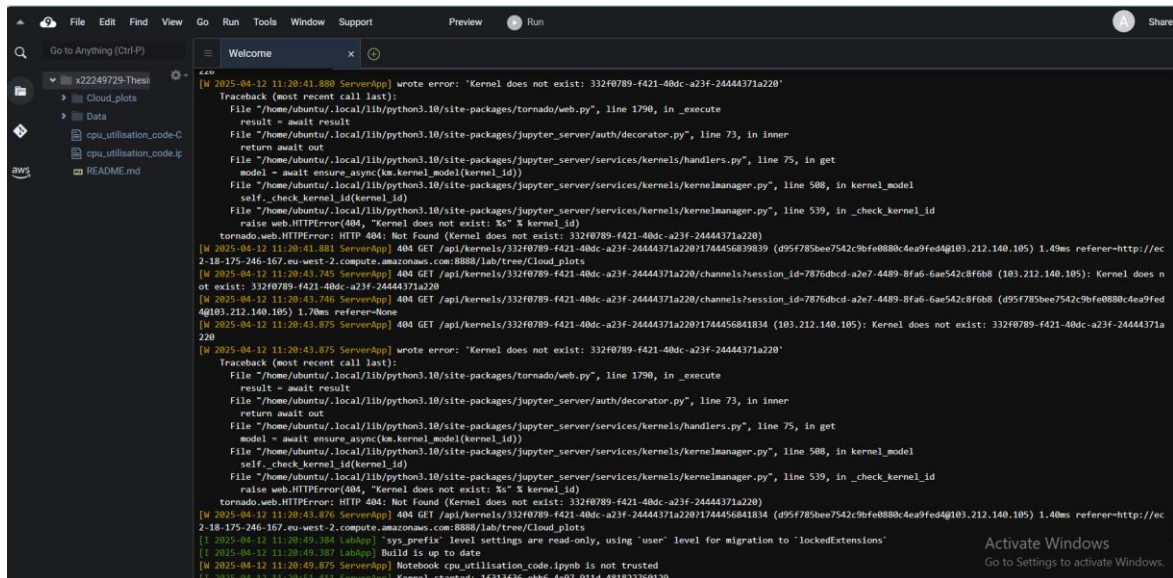
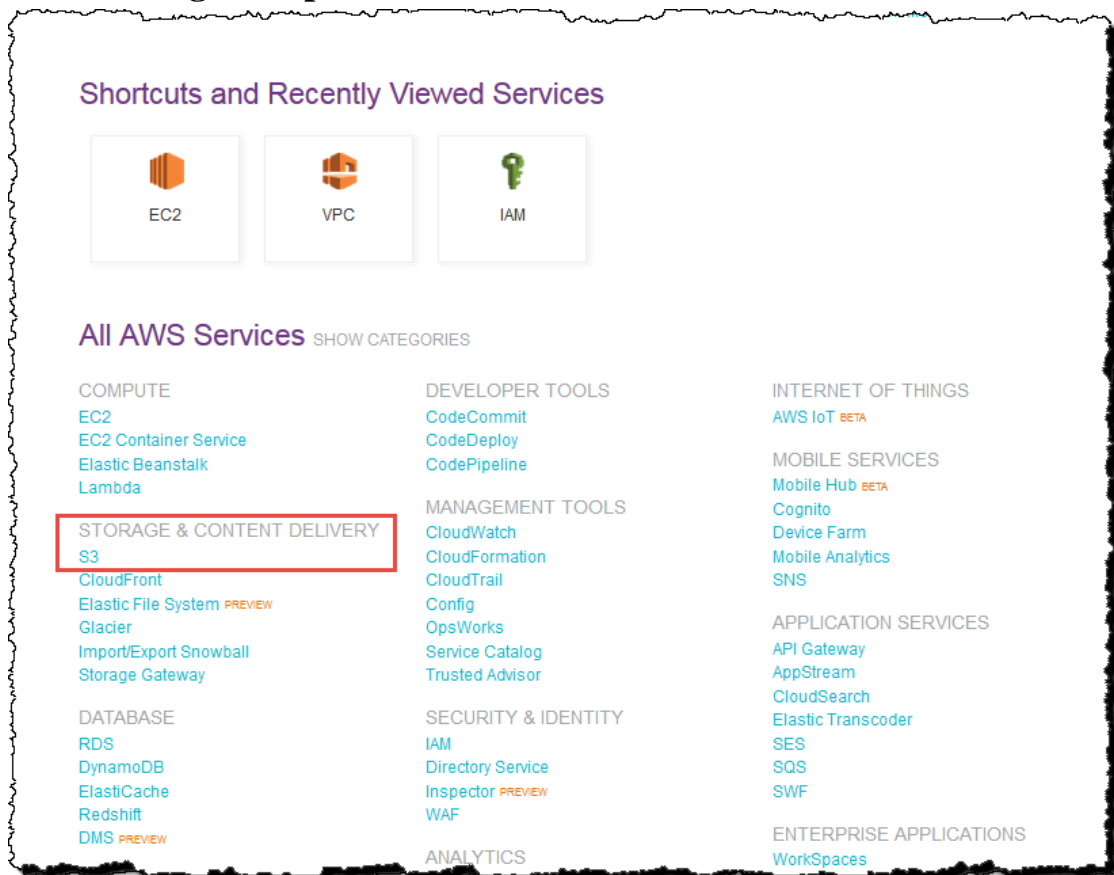


Figure 8: Running Jupyter from Cloud 9 IDE

## 2.4 S3 Storage Setup:



**Figure 9 : AWS Storage and Content Deliver under Console**

STEP 1: Find S3 under the Storage and content delivery menu in the AWS management console.



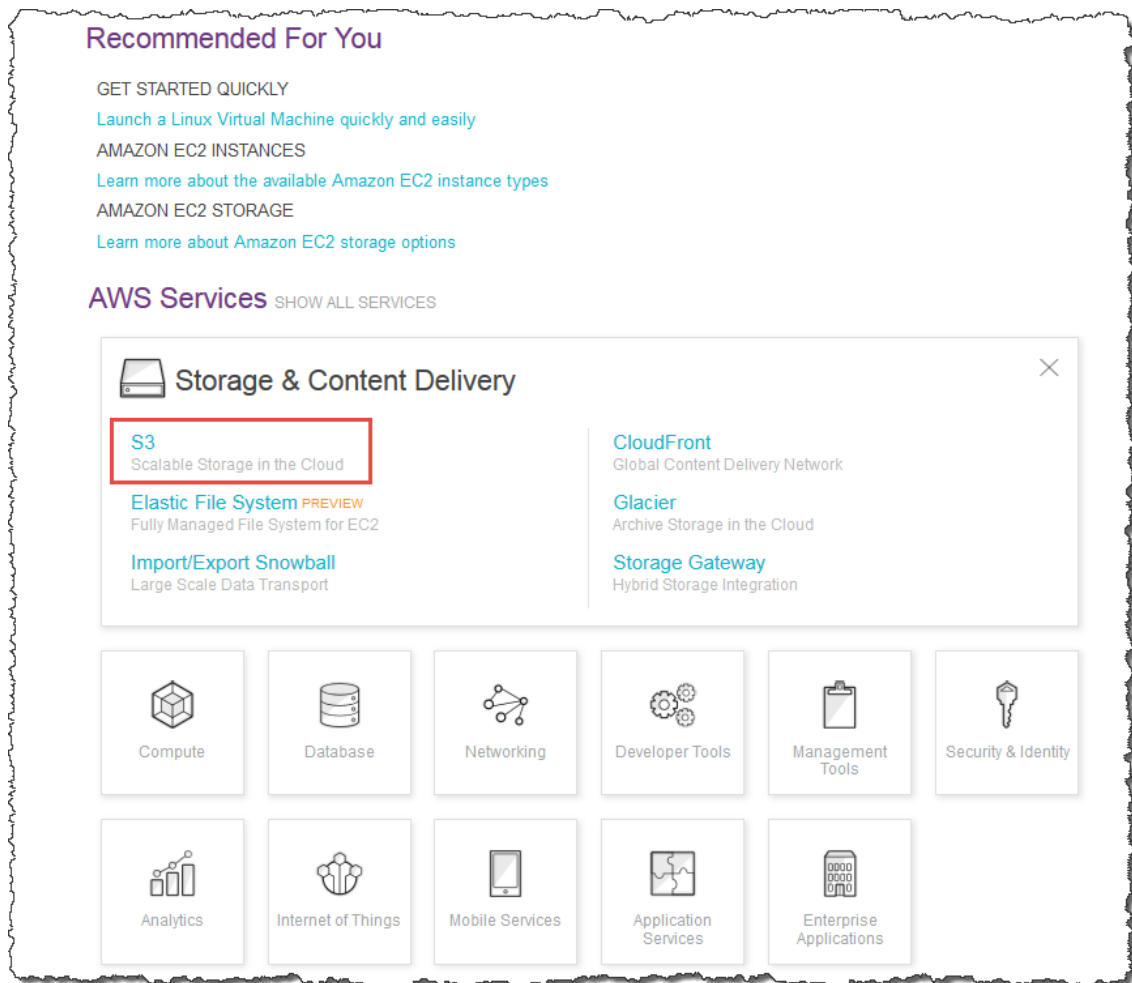


Figure 10: Select S3

STEP 2: Select S3

The image shows a screenshot of the 'Create a Bucket - Select a Bucket Name and Region' dialog box in the AWS console. The dialog has a title bar with a 'Cancel' button. The main content area contains a paragraph explaining that a bucket is a container for objects stored in Amazon S3 and that users can choose a region to optimize for latency, minimize costs, or address regulatory requirements. It also provides a link to 'Amazon S3 documentation'. Below the text, there are two input fields: 'Bucket Name:' with an empty text box, and 'Region:' with a dropdown menu currently showing 'Oregon'. At the bottom of the dialog, there are three buttons: 'Set Up Logging >', 'Create' (highlighted in blue), and 'Cancel'.

Figure 11 : Bucket Creation

STEP 3: Add the name of the bucket and select a region and click on create bucket.

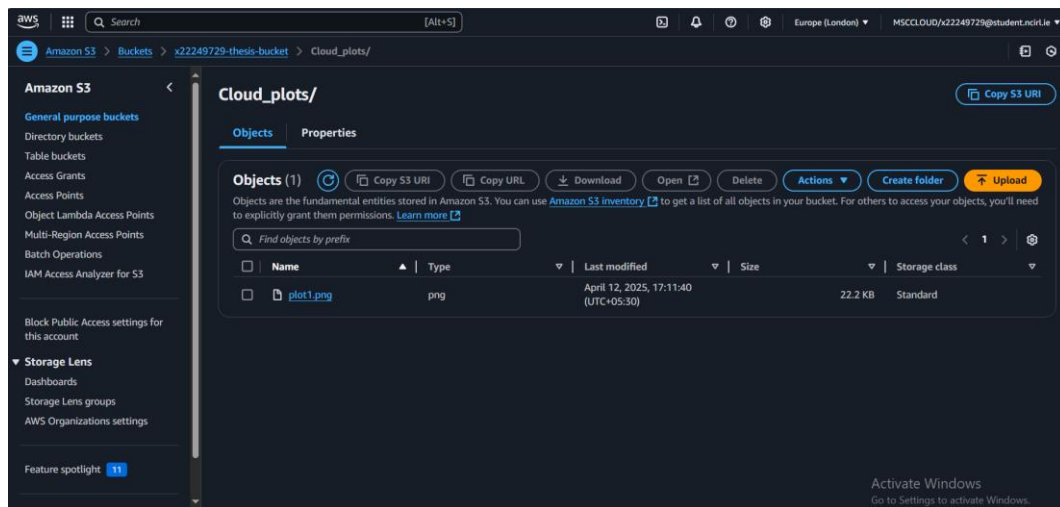


Figure 12: S3 bucket

## 3 Installing and Importing Python Libraries

### 3.1 Libraries and Packages

**Boto3:** An AWS SDK for Python used to interact with AWS services like S3 for cloud storage, enabling seamless data retrieval and management in cloud environments.

**Plotly (express and graph objects):** A powerful visualization library for creating interactive and customizable charts, graphs, and dashboards, perfect for exploring and presenting complex datasets.

**Seaborn:** A Python library for advanced statistical data visualization, offering aesthetic and informative charts to enhance exploratory data analysis (EDA).

**TensorFlow & Keras:** Components for building deep learning architectures such as LSTMs, Bidirectional LSTMs, and Cross-Attention mechanisms, essential for time-series and sequence modeling tasks.

**Scikit-learn:** Provides essential tools like Min-Max Scaler for data normalization and metrics like MSE and  $R^2$  for evaluating model performance.

### 3.2 Installing and Importing libraries

Open Jupyter Notebook from Cloud 9 IDE and install and import all the libraries as shown in Figure 13.

▼ Import All Libraries

```

]: import warnings
warnings.filterwarnings('ignore')

]: import glob
import boto3
import numpy as np
import pandas as pd
import seaborn as sns
import tensorflow as tf
import plotly.express as px
import plotly.offline as py
from sklearn import metrics
import matplotlib.pyplot as plt
import plotly.graph_objects as go
from tensorflow.keras import backend as K
from sklearn.preprocessing import MinMaxScaler
from botocore.exceptions import NoCredentialsError
from tensorflow.keras.models import Sequential, Model
from sklearn.model_selection import train_test_split
from tensorflow.keras.optimizers import Adam, SGD, RMSprop
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
from tensorflow.keras.layers import LSTM, Dense, Dropout, Bidirectional, Input, Permute, Multiply, Flatten

```

2024-12-07 06:36:27.387217: I external/local\_xla/xla/tsl/cuda/cudart\_stub.cc:32] Could not find cuda drivers on your machine, GPU will not be used.

2024-12-07 06:36:27.392376: I external/local\_xla/xla/tsl/cuda/cudart\_stub.cc:32] Could not find cuda drivers on your machine, GPU will not be used.

2024-12-07 06:36:27.405645: E external/local\_xla/xla/stream\_executor/cuda/cuda\_fft.cc:477] Unable to register cuFFT factory: Attempting to register factory for plugin cuFFT when one has already been registered

WARNING: All log messages before absl::InitializeLog() is called are written to STDERR

E0000 00:00:1733553387.428697 15160 cuda\_dnn.cc:8310] Unable to register cuDNN factory: Attempting to register factory for plugin cuDNN when one has already been registered

E0000 00:00:1733553387.435231 15160 cuda\_blas.cc:1418] Unable to register cuBLAS factory: Attempting to register factory for plugin cuBLAS when one has already been registered

2024-12-07 06:36:27.460695: I tensorflow/core/platform/cpu\_feature\_guard.cc:210] This TensorFlow binary is optimized to use available CPU instructions in

Figure 13: Importing Libraries

## 4 Implementation and Evaluation

### 4.1 Data Loading

#### Data Loading

[4]:

```

paths="materna_dataset/"

#List of all files
all_files = glob.glob(paths + "/*.csv")

datavalues = []
for filename in all_files[:150]: #Loading 150 file depending upon resource
    data = pd.read_csv(filename, delimiter=';', index_col=None, header=0)
    datavalues.append(data)

#concatenating the dataframe
dataframe = pd.concat(datavalues, axis=0, ignore_index=True)

```

[5]:

```
dataframe.head(10)
```

[5]:

	Timestamp	CPU cores	CPU capacity provisioned [MHZ]	CPU usage [MHZ]	CPU usage [%]	Memory capacity provisioned [KB]	Memory usage [KB]	Memory usage [%]	Disk read throughput [KB/s]	Disk write throughput [KB/s]	Disk size [GB]	Network received throughput [KB/s]	Network transmitted throughput [KB/s]
0	04.01.2016 00:00:00	2	0	115	2,51	4194304	122474	2,92	0	13	54	1	1
1	04.01.2016 00:05:00	2	0	106	2,32	4194304	119538	2,85	0	13	54	1	1
2	04.01.2016 00:10:00	2	0	119	2,61	4194304	329253	7,85	0	70	54	2	35
3	04.01.2016 00:15:00	2	0	121	2,63	4194304	284793	6,79	0	13	54	11	35

Figure 14: Loading 150 csv files using Pandas

### 4.2 Data Cleaning

Data Cleaning - Clean the data by addressing missing values, removing unnecessary columns, and fixing any other inconsistencies.

## Data Cleaning

```
[7]: dataframe.describe()
```

[7]:

	CPU cores	CPU capacity provisioned [MHZ]	CPU usage [MHZ]	Memory capacity provisioned [KB]	Memory usage [KB]	Disk read throughput [KB/s]	Disk write throughput [KB/s]	Disk size [GB]	Network received throughput [KB/s]	Network transmitted throughput [KB/s]
count	1.506811e+06	1506811.0	1.506811e+06	1.506811e+06	1.506811e+06	1.506811e+06	1.506811e+06	1.506811e+06	1.506811e+06	1.506811e+06
mean	3.203049e+00	0.0	3.751401e+02	1.035917e+07	9.046726e+05	6.424757e+02	1.370143e+02	4.847800e+02	1.489206e+02	1.949458e+02
std	1.946664e+00	0.0	8.787635e+02	6.361582e+06	1.339683e+06	7.278807e+03	1.063540e+03	2.874514e+03	9.816897e+02	1.564633e+03
min	1.000000e+00	0.0	1.600000e+00	2.097152e+06	0.000000e+00	0.000000e+00	0.000000e+00	5.400000e+01	0.000000e+00	0.000000e+00
25%	2.000000e+00	0.0	2.600000e+01	4.194304e+06	8.304700e+04	0.000000e+00	1.000000e+00	5.400000e+01	0.000000e+00	0.000000e+00
50%	2.000000e+00	0.0	7.300000e+01	8.388608e+06	3.456110e+05	0.000000e+00	5.000000e+00	1.080000e+02	1.000000e+00	1.000000e+00
75%	4.000000e+00	0.0	2.410000e+02	1.677722e+07	1.123235e+06	0.000000e+00	3.600000e+01	4.380000e+02	2.200000e+01	2.700000e+01
max	8.000000e+00	0.0	1.205600e+04	3.355443e+07	2.365252e+07	3.779170e+05	1.402520e+05	3.492800e+04	1.007140e+05	1.142190e+05

```
[8]: dataframe.info()
```

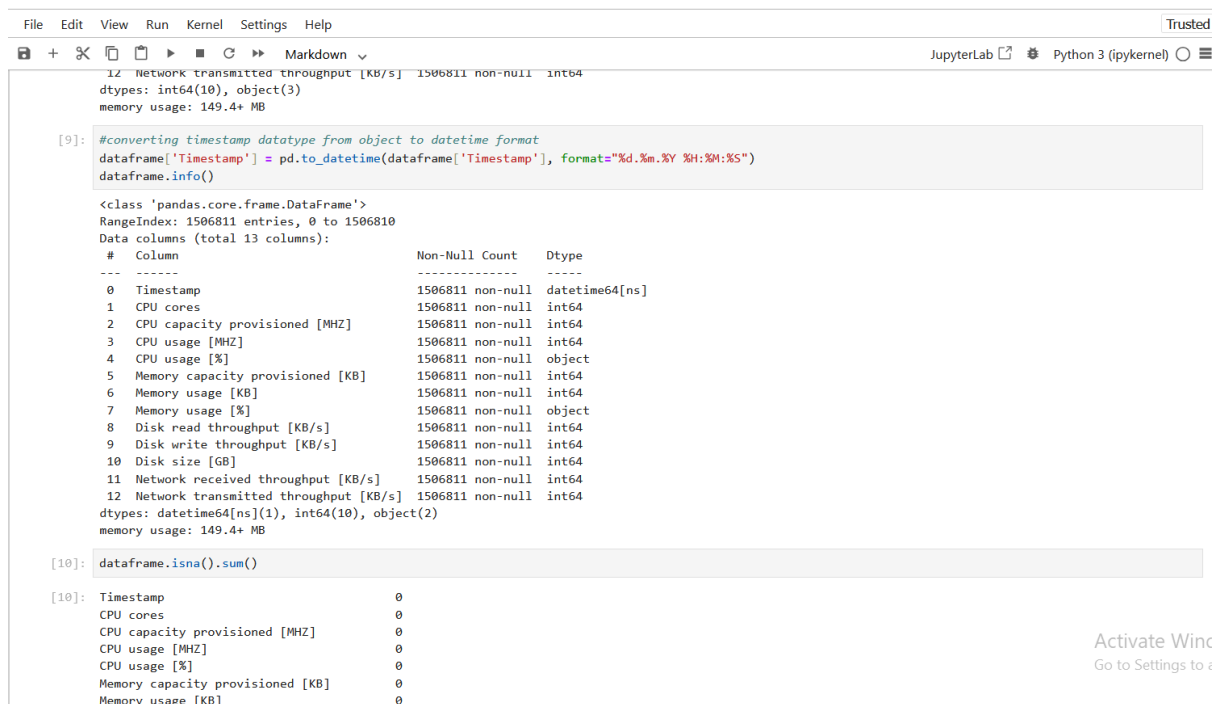
```

In [10]: df = pd.DataFrame(
RangeIndex: 1506811 entries, 0 to 1506810
Data columns (total 13 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Timestamp                             1506811 non-null  object
1   CPU cores                             1506811 non-null  int64
2   CPU capacity provisioned [MHZ]       1506811 non-null  int64
3   CPU usage [MHZ]                      1506811 non-null  int64

```

Activate Windows  
Go to Settings to activate Windows.

**Figure 15: Checking data inconsistency**



Activate Windows  
Go to Settings to activate Windows.

### Figure 16: Checking Null Values

### 4.3 Data Preprocessing and Feature Extraction

Process data and extract meaningful features which include datetime conversion of date columns, extracting month, year from date for visualization.

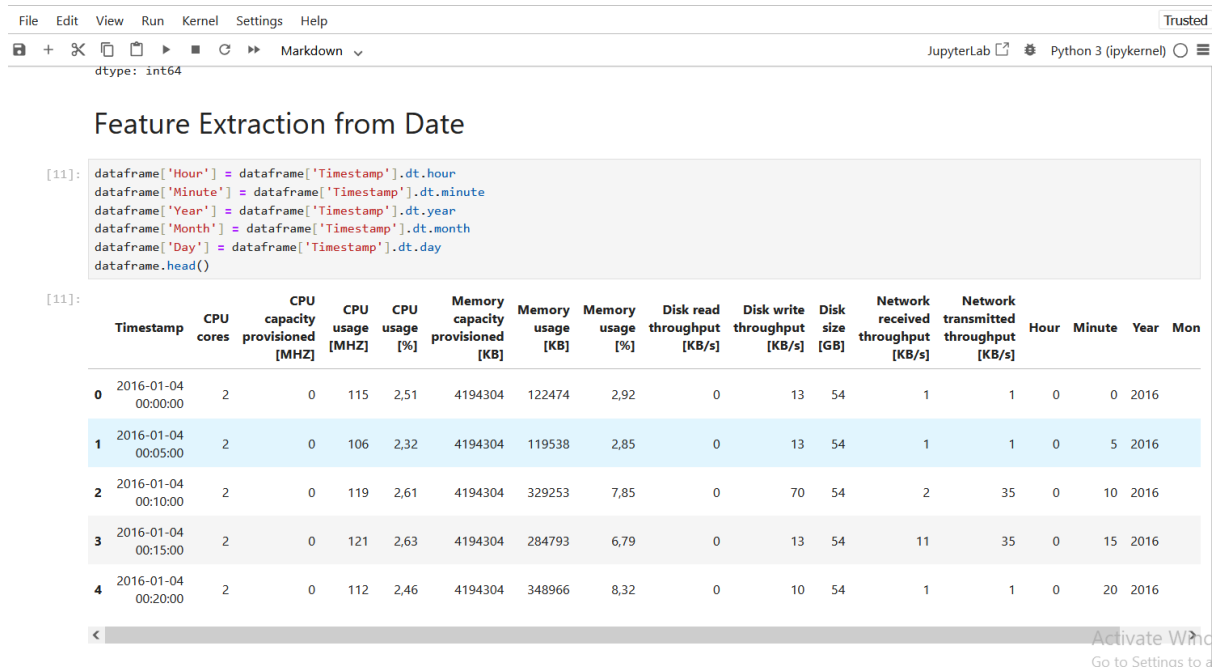


Figure 17: Feature extraction from Date

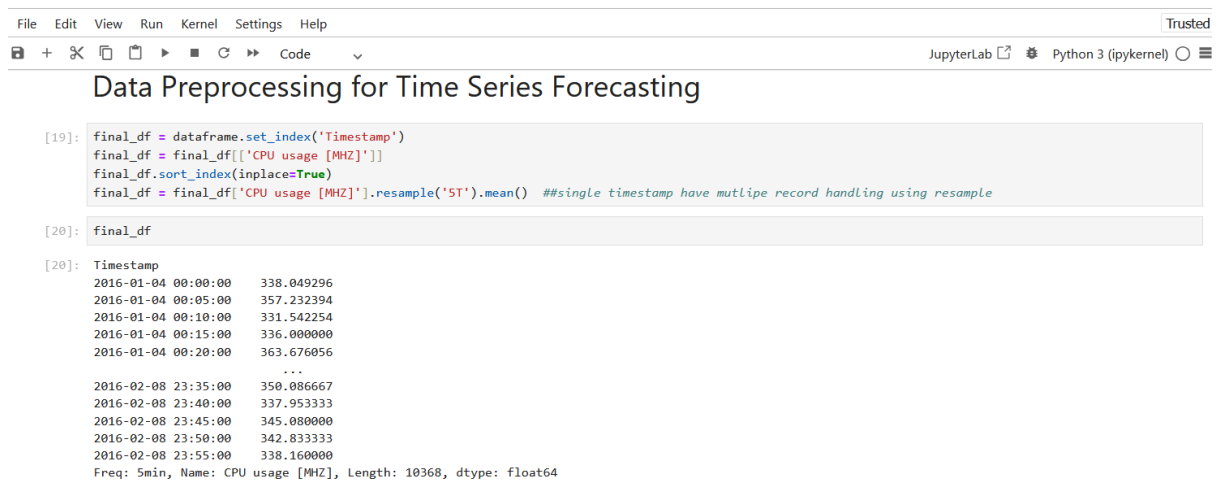




Figure 18: Preparing data for Forecasting

## 4.4 Data Visualization

Visualize the data to identify patterns, trends, and relationships between variables.




**Figure 19: Avg CPU Usage [MHz] by Month**




**Figure 20: Avg CPU Usage [MHz] by Day**


Avg CPU usage [MHz] by Hour



**Figure 21: Avg CPU Usage [MHz] by Hour**



**Figure 22: Avg CPU Usage [MHZ] by Min and Day**



**Figure 23: Time Series Plot of Usage by Date Time**

## 4.5 Scaling data – Min-Max Normalization

Normalize the data for improved model performance.

```
: #data normalization
data = final_df
scaler=MinMaxScaler(feature_range=(0,1))
dataframe=scaler.fit_transform(np.array(final_df).reshape(-1,1))
```

**Figure 24: Normalization of Data**

## 4.6 Window Rolling

Window Rolling with Timestamp (12 Previous Values)

```
: # apply window rolling timestep
def window_rolling(dataset, time_step=1):
    dataX, dataY = [], []
    for i in range(len(dataset)-time_step-1):
        a = dataset[i:(i+time_step), 0]
        dataX.append(a)
        dataY.append(dataset[i + time_step, 0])
    return np.array(dataX), np.array(dataY)

: time_step = 12
X_train, y_train = window_rolling(train_data, time_step)
X_test, y_test = window_rolling(test_data, time_step)
```

Figure 25: Apply window rolling to time step

```
[22]: #window rolling
new_final_df = final_df
rolling12_final_df = new_final_df.rolling(window=12).mean()
edadf = new_final_df.reset_index()
edadf1 = rolling12_final_df.reset_index()
newdf = edadf.merge(edadf1, on='Timestamp')
fig = px.line(newdf, x="Timestamp", y=newdf.columns, title='Actual vs WindowRolling Values')
fig.show()
file_name="./cpu_utilization_graphs/cloudcpuforecastgraph2.png"
fig.write_image(file_name)
upload_file_to_s3(file_name, bucket_name)
```




Figure 26: Comparing to actual values

## 4.7 Data Splitting

Split the data into training and testing (80:20) sets to evaluate model performance.

```
#splitting dataset into train and test with ration of 80:20
training_size=int(len(dataframe)*0.80)
test_size=len(dataframe)-training_size
train_data,test_data=dataframe[0:training_size,:],dataframe[training_size:len(dataframe),:1]
```

Figure 27: Splitting of Data



## 4.8 Data Learning Models

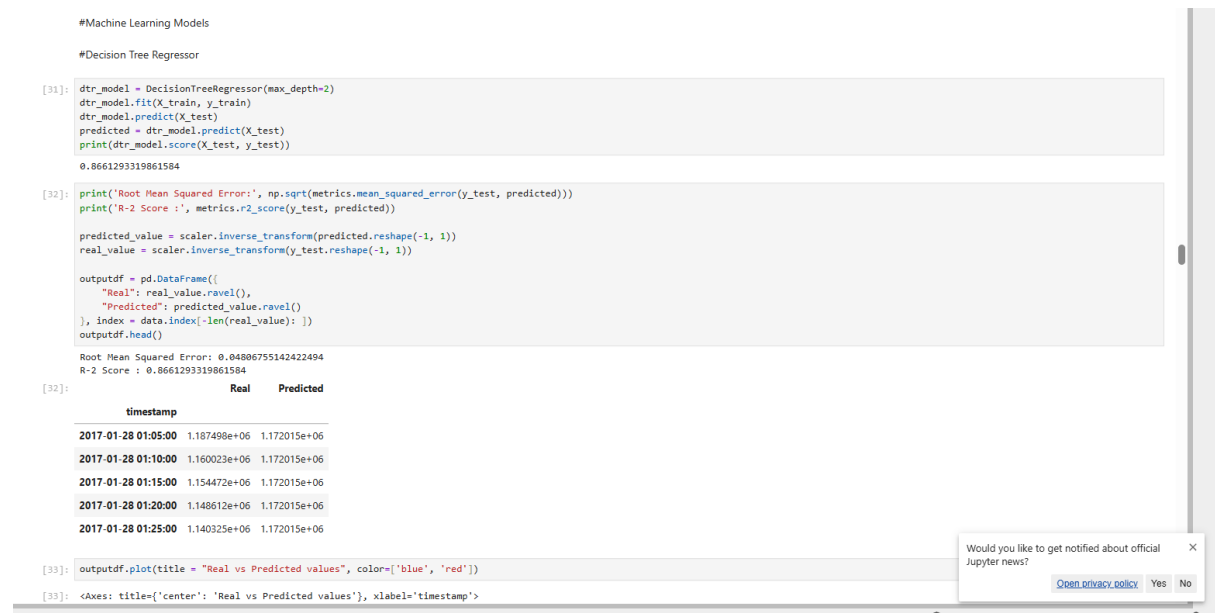


Figure 28: Running the Decision Tree Regressor

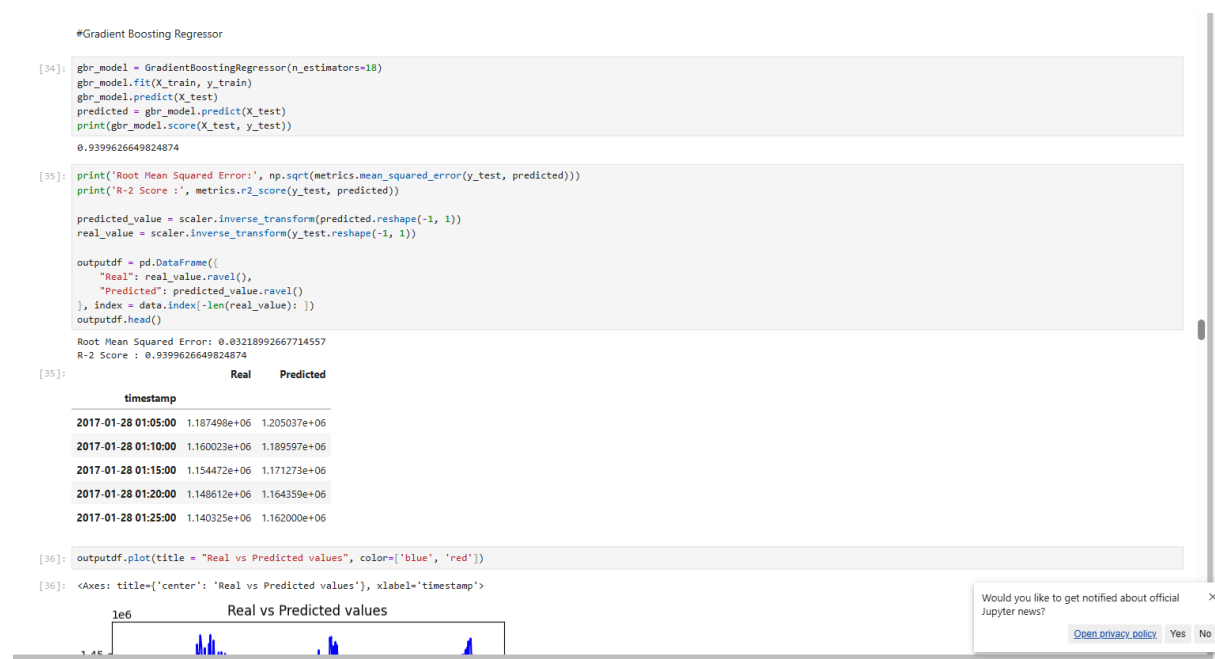


Figure 29: Running the Gradient Boosting Regressor

```
#LSTM Model

[38]: #lstm model
model = Sequential()
model.add(LSTM(units=12,return_sequences=True,input_shape=(X_train.shape[1], 1)))
model.add(Dropout(0.2))
model.add(LSTM(units=9, return_sequences=True))
model.add(Dropout(0.3))
model.add(LSTM(units=9))
model.add(Dropout(0.2))
model.add(Dense(1))

model.compile(optimizer="adam", loss="mean_squared_error")

[39]: model.summary()

Model: "sequential"

Layer (type)                 Output Shape              Param #
-----
lstm (LSTM)                   (None, 12, 12)            672
dropout (Dropout)             (None, 12, 12)            0
lstm_1 (LSTM)                 (None, 12, 9)             792
dropout_1 (Dropout)           (None, 12, 9)            0
lstm_2 (LSTM)                 (None, 9)                 684
dropout_2 (Dropout)           (None, 9)                 0
dense (Dense)                 (None, 1)                 10
-----
Total params: 2158 (8.43 KB)
Trainable params: 2158 (8.43 KB)
Non-trainable params: 0 (0.00 Byte)

[40]: plot_model(model, to_file='content/drive/MyDrive/cpu_utilisation_azure/model_plot_lstm.png', show_shapes=True, show_layer_names=True)

You must install pydot ('pip install pydot') and install graphviz (see instructions at https://graphviz.gitlab.io/download/) for plot_model to work.
```

Figure 30: Running Deep Learning Models (LSTM)

```
#BiLSTM Model

[48]: #biLstm model
model = Sequential()
model.add(Bidirectional(LSTM(units=32, input_shape=(X_train.shape[1], 1),return_sequences=True)))
model.add(Dropout(0.2))
model.add(LSTM(units=32, return_sequences=True))
model.add(Dropout(0.2))
model.add(LSTM(units=32))
model.add(Dropout(0.2))
model.add(Dense(units=32, kernel_initializer='uniform', activation='tanh'))
model.add(Dense(units=1, kernel_initializer='uniform', activation='linear'))

model.compile(optimizer='adam',loss='mean_squared_error')

[49]: history = model.fit(X_train,y_train,validation_data=(X_test,y_test),epochs=20,batch_size=128,verbose=1)

Epoch 1/20
61/61 [=====] - 9s 50ms/step - loss: 0.0292 - val_loss: 0.0042
Epoch 2/20
61/61 [=====] - 2s 27ms/step - loss: 0.0020 - val_loss: 0.0010
Epoch 3/20
61/61 [=====] - 2s 27ms/step - loss: 0.0018 - val_loss: 8.5279e-04
Epoch 4/20
61/61 [=====] - 2s 27ms/step - loss: 0.0016 - val_loss: 7.4491e-04
Epoch 5/20
61/61 [=====] - 2s 27ms/step - loss: 0.0015 - val_loss: 7.1637e-04
Epoch 6/20
61/61 [=====] - 2s 27ms/step - loss: 0.0014 - val_loss: 6.6889e-04
Epoch 7/20
61/61 [=====] - 2s 27ms/step - loss: 0.0013 - val_loss: 6.5247e-04
Epoch 8/20
61/61 [=====] - 2s 27ms/step - loss: 0.0012 - val_loss: 5.5515e-04
Epoch 9/20
61/61 [=====] - 2s 27ms/step - loss: 0.0012 - val_loss: 5.2999e-04
Epoch 10/20
61/61 [=====] - 2s 27ms/step - loss: 0.0011 - val_loss: 5.0553e-04
Epoch 11/20
61/61 [=====] - 2s 26ms/step - loss: 0.0010 - val_loss: 5.8510e-04
Epoch 12/20
61/61 [=====] - 2s 27ms/step - loss: 9.7958e-04 - val_loss: 5.0394e-04
Epoch 13/20
61/61 [=====] - 2s 27ms/step - loss: 0.0010 - val_loss: 6.7117e-04
Epoch 14/20
61/61 [=====] - 2s 27ms/step - loss: 9.3763e-04 - val_loss: 4.6932e-04
Epoch 15/20
61/61 [=====] - 2s 27ms/step - loss: 9.1375e-04 - val_loss: 4.9317e-04
```

Figure 31: BiLSTM Model Run

## Multihead bilstm

```
[56]: from tensorflow.keras.models import Model
from tensorflow.keras.layers import Input, LSTM, Bidirectional, Dense, Dropout, Concatenate

# Input Layer
input_layer = Input(shape=(X_train.shape[1], 1))

# === BiLSTM Head 1 ===
head1 = Bidirectional(LSTM(32, return_sequences=True))(input_layer)
head1 = Dropout(0.2)(head1)
head1 = LSTM(32, return_sequences=True)(head1)
head1 = Dropout(0.2)(head1)
head1 = LSTM(32)(head1)
head1 = Dropout(0.2)(head1)

# === BiLSTM Head 2 ===
head2 = Bidirectional(LSTM(32, return_sequences=True))(input_layer)
head2 = Dropout(0.2)(head2)
head2 = LSTM(32, return_sequences=True)(head2)
head2 = Dropout(0.2)(head2)
head2 = LSTM(32)(head2)
head2 = Dropout(0.2)(head2)

# === BiLSTM Head 3 (different configuration) ===
head3 = Bidirectional(LSTM(16, return_sequences=True))(input_layer)
head3 = Dropout(0.2)(head3)
head3 = LSTM(16)(head3)
head3 = Dropout(0.2)(head3)

# === Merge all heads ===
merged = Concatenate()([head1, head2, head3])

# === Dense Layers ===
dense = Dense(32, kernel_initializer='uniform', activation='tanh')(merged)
output = Dense(1, kernel_initializer='uniform', activation='linear')(dense)
```

Would you like to get notified about official Jupyter news?

[Open privacy policy](#) Yes

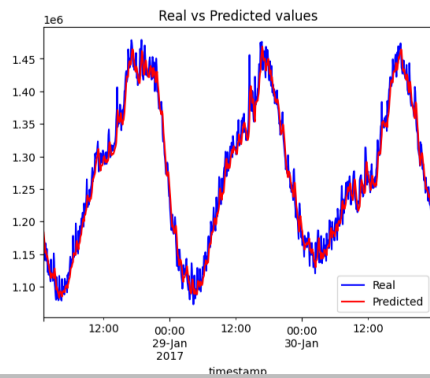
**Figure 32: Multihead BiLSTM Model Run**

```
outputter.head()
Root Mean Squared Error: 0.021343953440310356
R-Score : 0.9736044229354797

[60]:
```

	Real	Predicted
timestamp		
2017-01-28 01:05:00	1.187498e+06	1182782.875
2017-01-28 01:10:00	1.160023e+06	1184300.750
2017-01-28 01:15:00	1.154472e+06	1176124.125
2017-01-28 01:20:00	1.148612e+06	1168153.125
2017-01-28 01:25:00	1.140325e+06	1160207.000

```
[61]: outputdf.plot(title = "Real vs Predicted values", color=['blue', 'red'])
[61]: <Axes: title='center': 'Real vs Predicted values', xlabel='timestamp'>
```



Would you like to get notified about official Jupyter news?

[Open privacy policy](#) Yes No

**Figure 33: Model Results for Multihead BiLSTM Model**

## 5 Conclusion

The steps and procedures explained in this configuration manual will serve as a guide to any researcher trying to replicate the implementation of the cloud-based CPU Utilization forecasting system that is described in the research project. By following the instructions on the setup of the environment, data preprocessing and model training using all the AWS services used in this project by using the Python libraries, users can get the same experimental outcomes and performance metrics that have achieved in the current study. The visualizations and evaluation techniques demonstrate the objectives of the study that prove

the better performance of the deep learning models particularly Multihead BiLSTM in predicting time-series prediction in cloud environments.

## 6 References

Amazon Web Services (n.d.) *Step 1: Create an Amazon S3 bucket*. Available at: <https://docs.aws.amazon.com/quickstarts/latest/s3backup/step-1-create-bucket.html> (Accessed: 23 April 2025).

JFrog (n.d.) *Cloud9 setup – JFrog AWS Workshop*. Available at: [https://jfrog.awsworkshop.io/2\\_self\\_guided\\_setup/22\\_cloud9.html](https://jfrog.awsworkshop.io/2_self_guided_setup/22_cloud9.html) (Accessed: 23 April 2025).

Zhou, L. (n.d.) *Data visualization and exploration in pandas and matplotlib*. GitHub Gist. Available at: <https://gist.github.com/LiutongZhou/44655c8ed8e1c77d3f6a035e2b83e1f7> (Accessed: 23 April 2025).